

A Deterministic and Terminating Semantics for the Synchronous Programming Language CéU

Anonymous Author(s)

Abstract

CéU is a synchronous programming language for embedded soft real-time systems. It focus on control-flow safety features, such as safe shared-memory concurrency and safe abortion of lines of execution, while enforcing memory-bounded, deterministic, and terminating reactions to the environment. In this work, we present a small-step structural operational semantics for CéU and a proof that reactions have the properties enumerated above: that for a given arbitrary timeline of input events, multiple executions of the same program always react in bounded time and arrive at the same final finite memory state.

CCS Concepts • Theory of computation → Operational semantics; • Software and its engineering → Concurrent programming languages; • Computer systems organization → Embedded software;

Keywords Operational semantics, CéU, Synchronous languages

ACM Reference Format:

Anonymous Author(s). 1997. A Deterministic and Terminating Semantics for the Synchronous Programming Language CéU. In *Proceedings of ACM Woodstock conference (WOODSTOCK'97)*. ACM, New York, NY, USA, 10 pages. https://doi.org/10.475/123_4

1 Introduction

CéU [19, 21] is a Esterel-based [9] programming language for embedded soft real-time systems that aims to offer a concurrent, safe, and expressive alternative to C with the characteristics that follow:

Reactive: code only executes in reactions to events.

Structured: programs use structured control mechanisms, such as `await` (to suspend a line of execution), and `par` (to combine multiple lines of execution).

Synchronous: reactions run atomically and to completion on each line of execution, i.e., there's no implicit preemption or real parallelism.

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

WOODSTOCK'97, July 1997, El Paso, Texas USA

© 2016 Copyright held by the owner/author(s).

ACM ISBN 123-4567-24-567/08/06...\$15.00

https://doi.org/10.475/123_4

Structured reactive programming lets developers write code in direct style, recovering from the inversion of control imposed by event-driven execution [1, 16, 18]. Synchronous languages offer a simple run-to-completion execution model that enables deterministic execution and make formal reasoning tractable. For this reason, it has been successfully adopted in safety-critical real-time embedded systems [3].

Previous work in the context of embedded sensor networks evaluates the expressiveness of CéU in comparison to event-driven code in C and attests a reduction in source code size (around 25%) with a small increase in memory usage (around 5–10% for *text* and *data*) [21]. CéU has also been used in the context of multimedia systems [22] and games [20].

CéU inherits the synchronous and imperative mindset of Esterel but adopts a simpler semantics with fine-grained execution control [19]. The list that follows summarizes the semantic peculiarities of CéU:

- Fine-grained, intra-reaction deterministic execution, which makes CéU fully deterministic.
- Stack-based execution for internal events, which provides a limited but memory-bounded form of subroutines.
- Finalization mechanism for lines of execution, which makes abortion safe with regards to external resources.

In this work, we present a formal semantics for a subset of CéU that focus on its peculiarities in comparison to other synchronous languages.

- qual a abordagem / operational semantics / dois passos
- quais os resultados / provas
- quais os desafios e limitações
- **Guilherme:** TODO

Francisco: Descrever seções.

2 CéU

CéU is a synchronous reactive language in which programs advance in a sequence of discrete reactions to external events. It is designed for control-intensive applications, supporting concurrent lines of execution, known as *trails*, and instantaneous broadcast communication through events. Computations within a reaction (such as expressions, assignments, and system calls) are also instantaneous considering the synchronous hypothesis [10]. CéU provides an `await` statement that blocks the current running trail allowing the program to advance its other trails; when all trails are blocked, the reaction terminates and control returns to the environment.

In C  U, every execution path within loops must contain at least one `await` statement to an external input event [6, 21]. This restriction, which is statically checked by the compiler, ensures that every reaction runs in bounded time, eventually terminating with all trails blocked in `await` statements. C  U has an additional restriction, which it shares with Esterel and synchronous languages in general [4]: computations that take a non-negligible time to run (e.g., cryptography or image processing algorithms) violate the zero-delay hypothesis, and thus cannot be directly implemented.

Listing 1 shows a compact reference of C  U:

```

111 // Declarations
112 input <type> <id>;           // declares an external input event
113 event <type> <id>;           // declares an internal event
114 var <type> <id>;             // declares a variable
115
116 // Event handling
117 <id> = await <id>;           // awaits an event and assigns the received value
118 emit <id>(<exp>);           // emits an event passing a value
119
120 // Control flow
121 <stmt> ; <stmt>                // sequence
122 if <exp> then <stmts> else <stmts> end // conditional
123 loop do <stmts> end            // repetition
124 every <id> in <id> do <stmts> end    // event iteration
125 finalize [<stmts>] with <stmts> end // finalization
126
127 // Logical parallelism
128 par/or do <stmts> with <stmts> end // aborts when any side ends
129 par/and do <stmts> with <stmts> end // terminates when all sides ends
130 par do <stmts> with <stmts> end // never terminates
131
132 // Assignment & Integration with C
133 <id> = <exp>;                 // assigns a value to a variable
134 _<id>(<exp>);                 // calls a C function (id starts with '_')
```

Listing 1. The concrete syntax of C  U.

Listing 2 shows a complete example in C  U that toggles a LED whenever a radio packet is received, terminating with a button press always with the LED off. The implementation first declares the `BUTTON` and `RADIO_RECV` as input events (ln. 1–2). Then, it uses a `par/or` composition to run two activities in parallel: a single statement that waits for a button press before terminating (ln. 4), and an endless loop that toggles the LED on and off on radio receives (ln. 9–14). The `finalize` clause (ln. 6–8) ensures that, no matter how its enclosing trail terminates, the LED will be unconditionally turned off (ln. 7).

The `par/or` composition, which stands for a *parallel-or*, provides an orthogonal abortion mechanism [4] in which its composed trails do not know when and how they are aborted (i.e., abortion is external to them). The finalization mechanism extends orthogonal abortion to also work with activities that use stateful resources from the environment (such as files and network handlers), as we discuss in Section 2.3.

```

1 input void BUTTON;
2 input void RADIO_RECV;
3 par/or do
4   await BUTTON;
5 with
6   finalize with
7     _led(0);
8   end
9   loop do
10     _led(1);
11     await RADIO_RECV;
12     _led(0);
13     await RADIO_RECV;
14   end
15 end
```

Listing 2. A program in C  U that toggles a LED on every radio receive, terminating on a button press in a consistent state.

In C  U, any identifier prefixed with an underscore (e.g., `_led`) is passed unchanged to the underlying C compiler. Therefore, access to C is straightforward and syntactically traceable. To ensure that programs operate under the synchronous hypothesis, the compiler environment should only provide access to C operations that can be assumed to be instantaneous, such as non-blocking I/O and simple accesses to data structures.

2.1 External and Internal Events

C  U defines time as a discrete sequence of reactions to unique external input events. External input events are received from the environment, and each delimits a new logical unit of time that triggers an associated reaction. The life-cycle of a program in C  U can be summarized as follows [21]:

- i The program initiates a “boot reaction” in a single trail (parallel constructs may create new trails).
- ii Active trails execute until they await or terminate, one after another. This step is called a *reaction chain*, and always runs in bounded time.
- iii When all trails are blocked, the program goes idle and the environment takes control.
- iv On the occurrence of a new external input event, the environment awakes *all* trails awaiting that event, and the program goes back to step (ii).

A program must react to an event completely before handling the next one. By the synchronous hypothesis, the time the program spends in step (ii) is conceptually zero (in practice, negligible). Hence, from the point of view of the environment, the program is always idle on step (iii). In practice, if a new external input event occurs while a reaction executes, the event is saved on a queue, which effectively schedules it to be processed in a subsequent reaction.

External events and discrete time

The sequential processing of external input events induces a discrete notion of time in C  U, as illustrated in Figure 1. The continuous timeline shows an absolute reference clock with “physical timestamps” for the event occurrences (e.g., event C occurs at 17ms521us). The discrete timeline shows how the same occurring events fit in the logical notion of time of C  U. The boot reaction boot-   happens before any input, at program startup. Event A “physically” occurs during boot-   but, because time is discrete, its corresponding reaction only executes afterwards, at logical instant A-1. Similarly, event B occurs during A-1 and its reaction is postponed to execute at B-2. Event C also occurs during A-1 but its reaction must also wait for B-2 to execute and so it is postponed to execute at C-3. Event D occurs during an idle period and can start immediately at D-4. Finally, two instances of event E occur during D-4; they are handled in the subsequent reactions E-5 and E-6.

Unique input events imply mutually exclusive reactions, which execute atomically and never overlap. Automatic mutual exclusion is a prerequisite for deterministic reactions as we discuss in Section 3.

In practice, the synchronous hypothesis for C  U holds if reactions execute faster than the rate of incoming input events. Otherwise, the program would continuously accumulate delays between physical occurrences and actual reactions for the input events. Considering the context soft real-time systems, such delays and postponed reactions might be tolerated as long as they are infrequent and the application does not take too long to catch up with real time. Note that the synchronous semantics is the norm in typical event-driven systems, such as event dispatching in UI toolkits, game loops in game engines, and clock ticks in embedded systems.

Internal events as subroutines

In C  U, queue-based processing of events applies only to external input events, i.e., events submitted to the program by the environment. Internal events, which are events generated internally by the program via `emit` statements, are processed in a stack-based manner. Internal events provide a fine-grained execution control, and, because of their stack-based processing, can be used to implement a limited form of subroutines, as illustrated in Listing 3:




Figure 1. The discrete notion of time in C  U.

```

1 event int* inc;           // declares subroutine "inc"
2 par/or do
3   var int* p;
4   every p in inc do       // implements "inc" through an event iterator
5     *p = *p + 1;
6   end
7 with
8   var int v = 1;
9   emit inc(&v);           // calls "inc"
10  emit inc(&v);           // calls "inc"
11  _assert(v==3);          // asserts result after the two returns
12 end

```

Listing 3. A C  U program with a “subroutine”.

In the example, the “subroutine” `inc` is defined as an event iterator (ln. 4–6) that continuously awaits its identifying event (ln. 4), and increments the value passed by reference (ln. 5). A trail in parallel (ln. 8–11) invokes the subroutine through two consecutive `emit` statements (ln. 9–10). Given the stack-based execution for internal events, as the first `emit` executes, the calling trail pauses (ln. 9), the subroutine awakes (ln. 4), runs its body (yielding `v=2`), iterates, and awaits the next “call” (ln. 4, again). Only after this sequence does the calling trail resumes (ln. 9), makes a new invocation (ln. 10), and passes the assertion test (ln. 11).

C  U also supports nested `emit` invocations, e.g., the body of the subroutine `inc` (ln. 5) could `emit` an event targeting another subroutine, creating a new level in the stack. We can think of the stack as a record of the nested, fine-grained internal reactions that happen inside the same outer reaction to a single external event.

This form of subroutines has a significant limitation though: it cannot express recursion, since an `emit` to itself is always ignored as a running trail cannot be waiting on itself. That being said, it is this very limitation that brings important safety properties to subroutines. First, they are guaranteed to react in bounded time. Second, memory for locals is also bounded, not requiring data stacks.

At first sight, event iteration such as in “`every e do <...> end`” seems to be equivalent to “`loop do await e; <...> end`”. However, the loop variation would not compile because it does not contain a path to an external input `await` (`e` is an internal event). However, event iterators enforce syntactic restrictions and cannot contain `await` or `break` statements. The absence of `break` guarantees that an iterator never terminates from itself, essentially behaving as a safe blocking point in the program. For this reason, the restriction that execution paths within loops must contain at least one external `await` is extended to alternatively contain an `every` statement.

2.2 Shared-Memory Concurrency

Embedded applications make extensive use of global memory and shared resources, such as through memory-mapped

registers and system calls to device drivers. Hence, an important goal of C  U is to ensure a reliable behavior for programs with concurrent lines of execution sharing memory and interacting with the environment.

<pre> 336 input void A; 337 input void B; 338 var int x = 1; 339 par/and do 340 await A; 341 x = x + 1; 342 with 343 await B; 344 x = x * 2; 345 end </pre>	<pre> 1 input void A; 2 // (empty line) 3 var int y = 1; 4 par/and do 5 await A; 6 y = y + 1; 7 with 8 await A; 9 y = y * 2; 10 end </pre>
---	--

[a] Accesses to x are never concurrent. [b] Accesses to y are concurrent but deterministic.

Figure 2. Shared-memory concurrency in C  U: example [a] is safe because the trails access x atomically in different reactions; example [b] is unsafe because both trails access y in the same reaction.

In C  U, when multiple trails are active during the same reaction, they are scheduled in lexical order, i.e., in the order they appear in the program source code. For instance, consider the two examples in Figure 2, both defining shared variables (ln. 3), and assigning to them in parallel trails (ln. 6, 9).

In the example [a], the two assignments to x can only execute in reactions to different events A and B, which cannot occur simultaneously by definition (Section 2.1). Hence, for the sequence of events A→B, x becomes 4 ((1+1)*2), while for B→A, x becomes 3 ((1*2)+1).

In the example [b], the two assignments to y are simultaneous because they execute in reaction to the same event A. Since C  U employs lexical order for intra-reaction statements, the execution is still deterministic, and y always becomes 4 ((1+1)*2). However, note that an apparently innocuous change in the order of trails modifies the behavior of the program. To mitigate this threat, C  U performs concurrency checks at compile time to detect conflicting accesses to shared variables: if a variable is written in a trail segment, then a concurrent trail segment cannot read or write to that variable [21]. Nonetheless, the static checks are optional and are not a prerequisite for the deterministic semantics of the language.

2.3 Abortion and Finalization

The par/or of C  U is an orthogonal abortion mechanism because the two sides in the composition need not be tweaked with synchronization primitives or state variables in order to affect each other. In addition, abortion is *immediate* in the sense that it executes atomically in the current micro reaction.

<pre> par/or do var _msg_t msg; <...> // prepare msg finalize _send(&msg); with _cancel(&msg); end await SEND_ACK; with <...> end // </pre>	<pre> 1 par/or do 2 var _FILE* f; 3 finalize 4 f = _fopen(...); 5 with 6 _fclose(f); 7 end 8 _fwrite(..., f); 9 await A; 10 _fwrite(..., f); 11 with 12 <...> 13 end </pre>
---	---

[a] Local resource finalization [b] External resource finalization

Figure 3. C  U enforces the use of finalization to prevent *dangling pointers* for local resources and *memory leaks* for external resources.

Immediate orthogonal abortion is a distinctive feature of synchronous languages and cannot be expressed effectively in traditional (asynchronous) multi-threaded languages [4, 17].

However, aborting lines of execution that deal with external resources may lead to inconsistencies. For this reason, C  U provides a *finalize* construct to unconditionally execute a series of statements even if the enclosing block is externally aborted.

C  U also enforces the use of *finalize* for system calls that deal with pointers representing resources, as illustrated in the two examples of Figure 3:

- If C  U **passes** a pointer to a system call (ln. [a]:5), the pointer represents a **local** resource (ln. [a]:2) that requires finalization (ln. [a]:7).
- If C  U **receives** a pointer from a system call return (ln. [b]:4), the pointer represents an **external** resource (ln. [b]:2) that requires finalization (ln. [b]:6).

C  U tracks the interaction of system calls with pointers and requires finalization clauses to accompany them. In the example in Figure 3.a, the local variable msg (ln. 2) is an internal resource passed as a pointer to _send (ln. 5), which is an asynchronous call that transmits the buffer in the background. If the block aborts (ln. 11) before receiving an acknowledge from the environment (ln. 9), the local msg goes out of scope and the external transmission now holds a *dangling pointer*. The finalization ensures that the transmission also aborts (ln. 7). In the example in Figure 3.b, the call to _fopen (ln. 4) returns an external file resource as a pointer. If the block aborts (ln. 12) during the await A (ln. 9), the file remains open as a *memory leak*. The finalization ensures that the

file closes properly (ln. 6). In both cases, the code does not compile without the `finalize` construct.¹

The finalization mechanism of C    is fundamental to preserve the orthogonality of the `par/or` construct since the clean up code is encapsulated in the aborted trail itself.

3 Formal Semantics

In this section, we introduce a reduced syntax for C    and propose an operational semantics to formally describe the language. We describe a small synchronous kernel highlighting the peculiarities of C   , in particular the stack-based execution for internal events. For the sake of simplicity, we focus on the control aspects of the language, leaving out side effects and system calls (which behave like in conventional imperative languages).

3.1 Abstract Syntax

		<i>// primary expressions</i>
<code>p ::= mem(id)</code>		(any memory access to <code>'id'</code>)
<code>awaitExt(id)</code>		(await external event <code>'id'</code>)
<code>awaitInt(id)</code>		(await internal event <code>'id'</code>)
<code>emitInt(id)</code>		(emit internal event <code>'id'</code>)
<code>break</code>		(loop escape)
		<i>// compound expressions</i>
<code>if mem(id) then p else p</code>		(conditional)
<code>p ; p</code>		(sequence)
<code>loop p</code>		(repetition)
<code>every id p</code>		(event iteration)
<code>p and p</code>		(par/and)
<code>p or p</code>		(par/or)
<code>fin p</code>		(finalization)
		<i>// derived by semantic rules</i>
<code>p @loop p</code>		(unwinded loop)
<code>p @and q</code>		(unwinded par/and)
<code>p @or q</code>		(unwinded par/or)
<code>@canrun(n)</code>		(can run on stack level <code>'n'</code>)
<code>@nop</code>		(terminated expression)

Listing 4. Reduced syntax of C   .

Listing 4 shows the syntax for a subset of C    that is sufficient to describe all semantic peculiarities of the language. The `mem(id)` primitive represents all accesses, assignments, and system calls, that affect a memory location identified by `id`. According to the synchronous hypothesis of C   , `mem` expressions are considered to be atomic and instantaneous. As the challenging parts of C    reside on its control structures, we are not concerned here with a precise semantics for side effects, but only with their occurrences in programs. Note that `mem`, `awaitExt`, and `awaitInt/emitInt` expressions do not share identifiers, i.e., an identifier is either a variable, an external event, or an internal event.

Regarding mismatches between the concrete and formal languages, the concrete `await` and `emit` primitives support

¹ The compiler only forces the programmer to write the finalization clause, but cannot check if it actually handles the resource properly.

communication of values between them, e.g., an `emit a(10)` awakes a `v=await a` setting `v=10`. To reproduce this functionality in the formal semantics, we can use a shared variable to hold the value of an `emit` and access it after the corresponding `await`. Also, a `finalize <A> with end ; <C>` in the concrete language is equivalent to `<A> ; ((fin) or <C>)` in the abstract syntax. In the concrete language, `<A>` and `<C>` execute in sequence, and the finalization code `` is implicitly suspended waiting for `<C>` termination. In the formal language, `fin ` suspends forever when reached (it is an awaiting expression that never awakes). Hence, we need an explicit `or` to execute `<C>` in parallel, whose termination aborts `fin `, which finally causes `` to execute (by the semantic rules to be discussed).

3.2 Operational Semantics

The core of our semantics describes how a program reacts to a single external input event, i.e., starting from an input event, how the program behaves and becomes idle again to proceed to a subsequent reaction. We use a set of small-step operational rules, which are built in such a way that at most one transition is possible at any time, resulting in deterministic reactions. The transition rules map a triple with a program p , a stack level n , and an emitted event e to a modified triple as follows:

$$\langle p, n, e \rangle \rightarrow \langle p', n', e' \rangle \quad (\text{rule-general})$$

where

$$\begin{aligned} p, p' &\in P && (\text{program as described in Listing 4}) \\ n, n' &\in \mathbb{N} && (\text{current stack level}) \\ e, e' &\in \mathbb{E} \cup \{\epsilon\} && (\text{emitted event, possibly none}) \end{aligned}$$

At the beginning of a reaction to an input event `id`, the triple is initialized at stack level 0 ($n = 0$) and with the emitted event `id` ($e = id$). At the end of a reaction, after an arbitrary but finite number of transitions, the triple will block with a (possibly) modified program p' , at stack level 0, with no event emitted ϵ :

$$\langle p, 0, ext \rangle \xrightarrow{*} \langle p', 0, \epsilon \rangle$$

We also distinguish between two transition rule subtypes: *outermost* transitions \xrightarrow{out} ; and *nested* transitions \xrightarrow{nst} .

The \xrightarrow{out} rules **push** and **pop** are non-recursive definitions that only apply to the program as a whole and manipulate the stack level as follows:

$$\frac{e \neq \epsilon}{\langle p, n, e \rangle \xrightarrow{out} \langle bcast(p), n + 1, \epsilon \rangle} \quad (\text{push})$$

$$\frac{n > 0, ((p = @nop) \vee isblocked(n, p))}{\langle p, n, \epsilon \rangle \xrightarrow{out} \langle p, n - 1, \epsilon \rangle} \quad (\text{pop})$$

Rule **push** matches whenever there is an emitted event in the triple, and instantly broadcasts the event to the program, which means (a) awaking active wait expressions altogether (see $bcast(p)$ in Figure 4), (b) creating a nested reaction by increasing the stack level, and, at the same time, (c) consuming the event (e becomes ϵ). Note that rule **push** is the only one in the semantics that matches an emitted event and also immediately consumes it. The rule **pop** only decreases the stack level, not affecting the program, and only applies if the program is blocked (see $isblocked(n, p)$ in Figure 5). This condition ensures that an emit only resumes after its internal reaction completes and blocks, as discussed in Section 2.1.

Note that at the beginning of a reaction, an event ext is emitted, which will trigger rule **push**, which will immediately raise the stack level to 1. Since rule **pop** is the only to decrease the stack level and only applies to a blocked or terminated program, the semantics guarantees that a reaction only terminates with a blocked program at stack level 0.

The \xrightarrow{nst} rules are recursive definitions but do not affect the stack level and never have an emitted event as a precondition:

$$\langle p, n, \epsilon \rangle \xrightarrow{nst} \langle p', n, e \rangle \quad (\text{rule-inner})$$

The distinction between \xrightarrow{out} and \xrightarrow{nst} prevents rules **push** and **pop** to match and inadvertently modify the current stack level before completing a nested reaction. A complete reaction behaves as follows:

$$\begin{aligned} a) \quad & \langle p, 0, ext \rangle \xrightarrow[out]{push} \langle q, 1, \epsilon \rangle \\ b) \quad & [\xrightarrow[in]{*} \langle r, i, e \rangle \xrightarrow[out]{*} \langle s, j, \epsilon \rangle]^* \\ c) \quad & \xrightarrow[in]{*} \langle t, k, \epsilon \rangle \xrightarrow[out]{pop} \langle u, 0, \epsilon \rangle \end{aligned}$$

First (case a), a $\xrightarrow[out]{push}$ starts a nested reaction at level 1. Then (case b), a series of alternations between \xrightarrow{nst} transitions and a single \xrightarrow{out} transition represent, respectively, nested reactions and stack operations. Finally (case c), at some point, a $\xrightarrow[out]{pop}$ decrements the stack level to 0 and terminates the reaction.

The \xrightarrow{nst} transition rules for primary expressions are as follows:

$$\langle mem(id), n, \epsilon \rangle \xrightarrow{nst} \langle @nop, n, \epsilon \rangle \quad (\text{mem})$$

$$\langle emit(id), n, \epsilon \rangle \xrightarrow{nst} \langle @canrun(n), n, id \rangle \quad (\text{emitInt})$$

$$\langle @canrun(n), n, \epsilon \rangle \xrightarrow{nst} \langle @nop, n, \epsilon \rangle \quad (\text{canrun})$$

A mem operation becomes a $@nop$ which indicates the memory access (rule **mem**). An $emitInt(id)$ generates an event id and transits to a $@canrun(n)$ which can only resume at level n (rule **emitInt**). Since all \xrightarrow{nst} rules can only transit with $e = \epsilon$, an $emitInt$ causes rule **push** to execute at the outer level, creating a new level $n + 1$ on the stack. Also, with the new stack level, the resulting $@canrun(n)$ itself cannot transit, providing the desired stack-based semantics for internal events.

Proceeding to compound expressions, the rules for conditionals and sequences are straightforward:

$$\frac{val(id) \neq 0}{\langle (if \ mem(id) \ then \ p \ else \ q), n, \epsilon \rangle \xrightarrow{nst} \langle p, n, \epsilon \rangle} \quad (\text{if-true})$$

$$\frac{val(id, n) = 0}{\langle (if \ mem(id) \ then \ p \ else \ q), n, \epsilon \rangle \xrightarrow{nst} \langle q, n, \epsilon \rangle} \quad (\text{if-false})$$

$$\frac{\langle p, n, \epsilon \rangle \xrightarrow{nst} \langle p', n, e \rangle}{\langle (p ; q), n, \epsilon \rangle \xrightarrow{nst} \langle (p' ; q), n, e \rangle} \quad (\text{seq-adv})$$

$$\langle (@nop ; q), n, \epsilon \rangle \xrightarrow{nst} \langle q, n, \epsilon \rangle \quad (\text{seq-nop})$$

$$\langle (break ; q), n, \epsilon \rangle \xrightarrow{nst} \langle break, n, \epsilon \rangle \quad (\text{seq-brk})$$

Given that our semantics focuses on control, rules **if-true** and **if-false** are the only to query mem expressions. Function val receives a memory identifier and returns its current value.

The rules for loops are analogous to sequences, but use '@' as separators to properly bind breaks to their enclosing

loops:

$$\langle\langle loop\ p\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle p\ @loop\ p\rangle, n, \epsilon\rangle \quad (\text{loop-expd})$$

$$\frac{\langle p, n, \epsilon\rangle \xrightarrow{nst} \langle p', n, \epsilon\rangle}{\langle\langle p\ @loop\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle p'\ @loop\ q\rangle, n, \epsilon\rangle} \quad (\text{loop-adv})$$

$$\langle\langle @nop\ @loop\ p\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle loop\ p\rangle, n, \epsilon\rangle \quad (\text{loop-nop})$$

$$\langle\langle break\ @loop\ p\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle @nop\rangle, n, \epsilon\rangle \quad (\text{loop-brk})$$

When a program encounters a *loop*, it first expands its body in sequence with itself (rule **loop-expd**). Rules **loop-adv** and **loop-nop** are similar to rules **seq-adv** and **seq-nop**, advancing the loop until they reach a *@nop*. However, what follows the loop is the loop itself (rule **loop-nop**). Note that if we used ‘;’ as a separator in loops, rules **loop-brk** and **seq-brk** would conflict. Rule **loop-brk** escapes the enclosing loop, transforming everything into a *@nop*.

Proceeding to parallel compositions, the semantic rules for *and* and *or* always force transitions on their left branches *p* to occur before their right branches *q*:

$$\langle\langle p\ and\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle p\ @and\ (@canrun(n); q)\rangle, n, \epsilon\rangle \quad (\text{and-expd})$$

$$\frac{\langle p, n, \epsilon\rangle \xrightarrow{nst} \langle p', n, \epsilon\rangle}{\langle\langle p\ @and\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle p'\ @and\ q\rangle, n, \epsilon\rangle} \quad (\text{and-adv1})$$

$$\frac{isblocked(n, p), \langle q, n, \epsilon\rangle \xrightarrow{nst} \langle q', n, \epsilon\rangle}{\langle\langle p\ @and\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle p\ @and\ q'\rangle, n, \epsilon\rangle} \quad (\text{and-adv2})$$

$$\langle\langle p\ or\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle p\ @or\ (@canrun(n); q)\rangle, n, \epsilon\rangle \quad (\text{or-expd})$$

$$\frac{\langle p, n, \epsilon\rangle \xrightarrow{nst} \langle p', n, \epsilon\rangle}{\langle\langle p\ @or\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle p'\ @or\ q\rangle, n, \epsilon\rangle} \quad (\text{or-adv1})$$

$$\frac{isblocked(n, p), \langle q, n, \epsilon\rangle \xrightarrow{nst} \langle q', n, \epsilon\rangle}{\langle\langle p\ @or\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle p\ @or\ q'\rangle, n, \epsilon\rangle} \quad (\text{or-adv2})$$

Rules **and-expd** and **or-expd** insert a *@canrun(n)* at the beginning of the right branch. This ensures that an *emitInt* on the left branch, which transits to a *@canrun(n)*, still resumes before the right branch starts. The deterministic behavior of the semantics relies on the *isblocked* predicate (see Figure 5) which appears in rules **and-adv2** and **or-adv2**. These rules

require the left branch *p* to be blocked in order to allow the right branch transition from *q* to *q'*.

For a parallel *@and*, if one of the sides terminates, the composition is simply substituted by the other side (rules **and-nop1** and **and-nop2**). For a parallel *or*, if one of the sides terminates, the whole composition terminates, also applying the *clear* function to properly finalize the aborted side (rules **or-nop1** and **or-nop2**):

$$\langle\langle @nop\ @and\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle q, n, \epsilon\rangle \quad (\text{and-nop1})$$

$$\frac{isblocked(n, p)}{\langle\langle p\ @and\ @nop\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle p, n, \epsilon\rangle} \quad (\text{and-nop2})$$

$$\langle\langle @nop\ @or\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle clear(q), n, \epsilon\rangle \quad (\text{or-nop1})$$

$$\frac{isblocked(n, p)}{\langle\langle p\ @or\ @nop\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle clear(p), n, \epsilon\rangle} \quad (\text{or-nop2})$$

The *clear* function (see Figure 6) concatenates all active *fin* bodies of the side being aborted, so that they execute before the composition rejoins. Note that there are no transition rules for *fin* expressions. This is because once reached, a *fin* expression halts and will only execute when it is aborted by a trail in parallel and is expanded by the *clear* function. Note also that there is a syntactic restriction that *fin* bodies cannot contain awaiting expressions (*awaitExt*, *awaitInt*, *every*, and *fin*), i.e., they are guaranteed to execute entirely within a reaction.

Finally, a *break* in one of the sides in parallel escapes the closest enclosing *loop*, properly aborting the other side by applying the *clear* function:

$$\langle\langle break\ @and\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle clear(q); break\rangle, n, \epsilon\rangle \quad (\text{and-brk1})$$

$$\frac{isblocked(n, p)}{\langle\langle p\ @and\ break\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle clear(p); break\rangle, n, \epsilon\rangle} \quad (\text{and-brk2})$$

$$\langle\langle break\ @or\ q\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle clear(q); break\rangle, n, \epsilon\rangle \quad (\text{or-brk1})$$

$$\frac{isblocked(n, p)}{\langle\langle p\ @or\ break\rangle, n, \epsilon\rangle \xrightarrow{nst} \langle\langle clear(p); break\rangle, n, \epsilon\rangle} \quad (\text{or-brk2})$$

A reaction eventually blocks in *awaitExt*, *awaitInt*, *every*, *fin*, and *@canrun* expressions in parallel trails. If no trails are blocked in *@canrun* expressions, it means that the program cannot advance in the current reaction. However, *@canrun* expressions can still resume in lower stack indexes and will eventually resume in the current reaction (see rule **pop**).

```

771  bcast(e, awaitExt(e)) = @nop
772  bcast(e, awaitInt(e)) = @nop
773
774  bcast(e, every e p) = p; every e p
775  bcast(e, @canrun(n)) = @canrun(n)
776  bcast(e, fin p) = fin p
777  bcast(e, p ; q) = bcast(e, p) ; q
778  bcast(e, p @loop q) = bcast(e, p) @loop q
779  bcast(e, p @and q) = bcast(e, p) @and bcast(e, q)
780  bcast(e, p @or q) = bcast(e, p) @or bcast(e, q)
781  bcast(e, _) = ⊥ (mem, emitInt, break, if,
782                    loop, and, or, @nop)
783
784
785

```

Figure 4. The function *bcast* awakes awaiting trails matching the event by converting *awaitExt* and *awaitInt* to *@nop* expressions, and by unwinding *every* expressions.

```

791  isblocked(n, awaitExt(id)) = true
792  isblocked(n, awaitInt(id)) = true
793  isblocked(n, every e p) = true
794  isblocked(n, @canrun(m)) = (n > m)
795  isblocked(n, fin p) = true
796  isblocked(n, p ; q) = isblocked(n, p)
797  isblocked(n, p @loop q) = isblocked(n, p)
798  isblocked(n, p @and q) = isblocked(n, p) ∧ isblocked(n, q)
799  isblocked(n, p @or q) = isblocked(n, p) ∧ isblocked(n, q)
800  isblocked(n, _) = false (mem, emitInt, break, if,
801                          loop, and, or, @nop)
802
803
804
805

```

Figure 5. The predicate *isblocked* is true only if all branches in parallel are blocked waiting for events, finalization clauses, or certain stack levels.

3.3 Determinism, Termination, and Memory Bounds

- informal discussion

Determinism

The proof for determinism relies on the fact all semantic rules are mutually exclusive, i.e., their preconditions are unique in the set of rules. This can be verified by direct inspection of rules.

Rule **push** is the only one with $e \neq \epsilon$ as a precondition, and is trivially mutually exclusive with all other rules.

Rule **pop** either has $p = @nop$ or *isblocked*(n, p) as preconditions. Note that rule **pop** only applies syntactically to top-level transitions. For instance, it can never match \xrightarrow{nst}

rules for subprograms as in rule **seq-adv**. Hence, for the first case, rule **pop** only applies, and is the only one to apply, to *nop* as the whole program (i.e., a *nop* not surrounded by other expressions, such as in rule **seq-nop**). For the second case, we need to show that given $\langle p, n, \epsilon \rangle$, no \xrightarrow{nst} transitions apply with *isblocked*(n, p) and vice versa. Except for *@canrun*, there are no \xrightarrow{nst} transitions for the other blocking expressions (*awaitExt*, *awaitInt*, *every*, and *fin*). However, considering the precondition $\langle p, n, \epsilon \rangle$, *isblocked*($n, @canrun(n)$) is false. Hence, given the preconditions for rule **pop**, no \xrightarrow{nst} transitions can occur. Conversely, if a \xrightarrow{nst} transition is possible, then *isblocked*(n, p) must be false. Again, except for *@canrun*, all other transitions do not involve blocking expressions, hence, for these transitions, *isblocked*(n, p) must be false. For rule **canrun**, a transition can only occur if the current stack level matches *@canrun*(n). In this case, *isblocked*($n, @canrun(n)$) is false.

Finally, we need to show that \xrightarrow{nst} transitions are mutually exclusive among themselves. Note that most rules have unique syntactic prefixes, e.g., (*@nop @and q*) (rule **and-nop1**) is trivially mutually exclusive with (*@nop @loop p*) (rule **or-nop1**). The only exceptions are rules **and-adv1** vs. **and-adv2**, and **or-adv1** vs. **or-adv2**. In both cases, we need to show that if the left branch can advance, then it cannot be blocked and vice-versa, i.e., that $\langle p, n, \epsilon \rangle \xrightarrow{nst} \langle p', n, \epsilon \rangle$ and *isblocked*(n, p) are mutually exclusive, which is exactly the same reasoning for rule **pop** above.

Termination

- there is always a possible transition until $n=0$
- every cannot restart itself - break disallowed - emit ignored

```

826  clear(awaitExt(e)) = @nop
827  clear(awaitInt(e)) = @nop
828
829  clear(every e p) = @nop
830  clear(@canrun(n)) = @nop
831  clear(fin p) = p
832  clear(p ; q) = clear(p)
833  clear(p @loop q) = clear(p)
834  clear(p @and q) = clear(p) ; clear(q)
835  clear(p @or q) = clear(p) ; clear(q)
836  clear(_) = ⊥ (mem, emitInt, break, if,
837                loop, and, or, @nop)
838
839
840

```

Figure 6. The function *clear* extracts *fin* expressions in parallel and put their bodies in sequence.

Memory Bounds

- program is finite - lexical scope - no heap allocation - no code reentrancy - reexecution only due to loops - loop reuse nested vars -

4 Related Work

C  u was strongly influenced by Esterel but they differ in the most fundamental aspect of the notion of time [21]. Esterel defines time as a discrete sequence of logical unit instants or “ticks”. At each tick, the program reacts to an arbitrary number of simultaneous input events from the environment. The presence of multiple inputs requires careful static analysis to detect and reject programs with *causality cycles* and *schizophrenia problems* [5, 8, 12, 23–26, 30]. In contrast, C  u defines time as a discrete sequence of reactions to unique input events. In the formal semantics, ...

Francisco: como isso aparece na semantica

... C  u also rejects some syntactically correct programs to avoid infinite execution, but with simple restrictions in the abstract syntax tree.

Another distinction is that, in Esterel, the behavior of internal and external events is equivalent, while in C  u internal events introduce stack-based micro reactions within external reactions, providing more fine-grained control for intra-reaction execution. In the formal semantics, ...

Francisco: como isso aparece na semantica

... Some variants of the Statecharts synchronous visual language also distinguish internal from external events [28]. In Statemate [14], “*reactions to external and internal events (...) can be sensed only after completion of the step*”, implying queue-based execution. In Stateflow [13], “*the receiving state (of the event) acts here as a function*”, which is similar to C  u’s stack-based execution. We are not aware of formalizations for these ideas for a deeper comparison with C  u.

Like other synchronous languages (*Reactive C* [7], *Protothreads* [11], *SOL* [15], *SC* [29], and *PRET-C* [2]), C  u relies on deterministic scheduling to preserve intra-reaction determinism. In addition, it also performs concurrency checks to detect trails that, when reordered, change the observable behavior of the program, i.e., trails that actually rely on deterministic scheduling [21]. Esterel is only deterministic with respect to external behavior: “*the same sequence of inputs always produces the same sequence of outputs*” [6]. However, the execution order for operations within a reaction is non-deterministic: “*if there is no control dependency, as in (call f1() || call f2()), the order is unspecified and it would be an error to rely on it*” [6]. For this reason, Esterel, does not support shared-memory concurrency: “*if a variable is written by some thread, then it can neither be read nor be written by concurrent threads*” [6].

Esterel describes a finalization mechanism in a standardization proposal [27] that is similar to C  u’s. However, we are

not aware of an open implementation or a formal semantics for a deeper comparison.

Francisco: outras linguagens sincronas

Francisco: outras linguagens deterministicas

Francisco: outras linguagens com terminacao

5 Conclusion

Francisco: TODO

References

- [1] A. Adya et al. 2002. Cooperative Task Management Without Manual Stack Management. In *Proceedings of ATEC’02*. USENIX Association, 289–302.
- [2] Sidharta Andalam, Partha Roop, and Alain Girault. 2010. Predictable multithreading of embedded applications using PRET-C. In *Proceeding of MEMOCODE’10*. IEEE, 159–168.
- [3] Albert Benveniste, Paul Caspi, Stephen A. Edwards, Nicolas Halbwachs, Paul Le Guernic, and Robert De Simone. 2003. The synchronous languages twelve years later. In *Proceedings of the IEEE*, Vol. 91. 64–83.
- [4] G  rard Berry. 1993. Preemption in Concurrent Systems.. In *FSTTCS (LNCS)*, Vol. 761. Springer, 72–93.
- [5] G  rard Berry. 1999. *The Constructive Semantics of Pure Esterel (draft version 3)*. Ecole des Mines de Paris and INRIA.
- [6] G  rard Berry. 2000. *The Esterel-V5 Language Primer*. CMA and Inria, Sophia-Antipolis, France. Version 5.10, Release 2.0.
- [7] Fr  d  ric Boussinot. 1991. Reactive C: An extension of C to program reactive systems. *Software: Practice and Experience* 21, 4 (1991), 401–428.
- [8] Fr  d  ric Boussinot. 1998. SugarCubes implementation of causality. (1998).
- [9] Fr  d  ric Boussinot and Robert De Simone. 1991. The Esterel language. *Proc. IEEE* 79, 9 (Sep 1991), 1293–1304.
- [10] Robert de Simone, Jean-Pierre Talpin, and Dumitru Potop-Butucaru. 2005. The Synchronous Hypothesis and Synchronous Languages. In *Embedded Systems Handbook*, R. Zurawski (Ed.).
- [11] Adam Dunkels, Oliver Schmidt, Thimo Voigt, and Muneeb Ali. 2006. Protothreads: simplifying event-driven programming of memory-constrained embedded systems. In *Proceedings of SenSys’06*. ACM, 29–42.
- [12] Stephen A. Edwards. 2005. Using and Compiling Esterel. MEMOCODE’05 Tutorial. (July 2005).
- [13] Gr  goire Hamon and John Rushby. 2007. An operational semantics for Stateflow. *International Journal on Software Tools for Technology Transfer* 9, 5-6 (2007), 447–456.
- [14] David Harel and Amnon Naamad. 1996. The STATEMATE semantics of statecharts. *ACM Transactions on Software Engineering and Methodology* 5, 4 (1996), 293–333.
- [15] Marcin Karpinski and Vinny Cahill. 2007. High-Level Application Development is Realistic for Wireless Sensor Networks. In *Proceedings of SECON’07*. 610–619.
- [16] Ingo Maier, Tiark Rompf, and Martin Odersky. 2010. *Deprecating the observer pattern*. Technical Report.
- [17] ORACLE. 2011. Java Thread Primitive Deprecation. <http://docs.oracle.com/javase/6/docs/technotes/guides/concurrency/threadPrimitiveDeprecation.html> (accessed in Aug-2014). (2011).
- [18] Guido Salvaneschi et al. 2014. REScala: Bridging between object-oriented and functional style in reactive applications. In *Proceedings of Modularity’13*. ACM, 25–36.
- [19] Francisco Sant’anna, Roberto Ierusalimsky, Noemi Rodriguez, Silvana Rossetto, and Adriano Branco. 2017. The Design and Implementation

- of the Synchronous Language CÉU. *ACM Trans. Embed. Comput. Syst.* 16, 4, Article 98 (July 2017), 26 pages. <https://doi.org/10.1145/3035544>
- [20] Francisco Sant'Anna, Noemi Rodriguez, and Roberto Ierusalimschy. 2015. Structured Synchronous Reactive Programming with Céu. In *Proceedings of Modularity'15*.
- [21] Francisco Sant'Anna, Noemi Rodriguez, Roberto Ierusalimschy, Olaf Landsiedel, and Philippas Tsigas. 2013. Safe System-level Concurrency on Resource-Constrained Nodes. In *Proceedings of SenSys'13*. ACM.
- [22] Rodrigo Santos, Guilherme Lima, Francisco Sant'Anna, and Noemi Rodriguez. 2016. Céu-Media: Local Inter-Media Synchronization Using Céu. In *Proceedings of WebMedia'16*. ACM, New York, NY, USA, 143–150. <https://doi.org/10.1145/2976796.2976856>
- [23] Klaus Schneider and Michael Wenz. 2001. A new method for compiling schizophrenic synchronous programs. In *Proceedings of CASES'01*. ACM, 49–58.
- [24] Ellen M Sentovich. 1997. Quick conservative causality analysis. In *System Synthesis, 1997. Proceedings., Tenth International Symposium on*. IEEE, 2–8.
- [25] Thomas R Shiple, Gerard Berry, and Hemé Touati. 1996. Constructive analysis of cyclic circuits. In *European Design and Test Conference, 1996. ED&TC 96. Proceedings*. IEEE, 328–333.
- [26] Olivier Tardieu and Robert De Simone. 2004. Curing schizophrenia by program rewriting in Esterel. In *Proceedings of MEMOCODE'04*. IEEE, 39–48.
- [27] Esterel Technologies. 2005. *The Esterel v7 Reference Manual*. Initial IEEE standardization proposal.
- [28] Michael von der Beeck. 1994. A comparison of statecharts variants. In *Proceedings of FTRTFT'94*. Springer, 128–148.
- [29] Reinhard von Hanxleden. 2009. SyncCharts in C: a proposal for lightweight, deterministic concurrency. In *Proceedings EMSOFT'09*. ACM, 225–234.
- [30] Jeong-Han Yun, Chul-Joo Kim, Seonggun Kim, Kwang-Moo Choe, and Taisook Han. 2013. Detection of harmful schizophrenic statements in Esterel. *ACM Transactions on Embedded Computing Systems (TECS)* 12, 3 (2013), 80.

A Proofs